


```
In [1]: from sklearn.linear_model import LinearRegression
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import linear_model
```

```
In [2]: n = 10000
x = np.linspace(0.01, 1, n).reshape(-1, 1)
y = np.linspace(0.01, 1, n) + np.random.rand(n) - .5

plt.scatter(x,y)
```

Out[2]: <matplotlib.collections.PathCollection at 0x22dca7cda20>



Assignment 5

1. Create and fit a Linear Regression Model

Calculate the Training error and Testing error using sklearn with a .50 split

For error, use mean_squared, but if you want to experiment with other mean errors, please do!

```
In [37]: from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.5)

lr = LinearRegression()
lr.fit(xtrain,ytrain)

train_err = mean_squared_error(ytrain, np.dot(xtrain, lr.coef_) + lr.intercept_)
test_err = mean_squared_error(ytest, np.dot(xtest,lr.coef_) + lr.intercept_)

print('train error: {0}'.format(train_err))
print('test error: {0}'.format(test_err))
print('slope:{0}'.format(lr.coef_))

train error: 0.08447252784478264
test error: 0.08368899099673137
slope:[1.02132297]
```

2. Repeat #1 for a Ridge Regression

```
In [4]: from sklearn.linear_model import Ridge
ridge = Ridge(fit_intercept=True)
ridge.fit(xtrain, ytrain)

train_err = mean_squared_error(ytrain, np.dot(xtrain, ridge.coef_) + ridge.int
ercept_)
test_err = mean_squared_error(ytest, np.dot(xtest,ridge.coef_) + ridge.interc
ept_)

print('train error: {0}'.format(train_err))
print('test error: {0}'.format(test_err))

train error: 0.08402671905858511
test error: 0.08417675172400539
```

3. Vary the split size from .01 to .99 with at least 10 values (the more the merrier!). Plot the resulting Training error and Testing error vs. split size. Create separate plots for Linear and Ridge

```

In [29]: import sys
splits = np.linspace(.01, .99,20)
np.random.seed(1256)
errValuesRegression = []
errValuesRidge = []

min_err = (sys.float_info.max, 0.0) #min err, split

for split in splits:
    xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=split)

    lr = LinearRegression()
    lr.fit(xtrain,ytrain)

    train_err = mean_squared_error(ytrain, np.dot(xtrain, lr.coef_) + lr.intercept_)
    test_err = mean_squared_error(ytest, np.dot(xtest,lr.coef_) + lr.intercept_)

    errValuesRegression.append((train_err, test_err))

    ridge = Ridge(alpha=100.0,fit_intercept=True)
    ridge.fit(xtrain, ytrain)

    train_err2 = mean_squared_error(ytrain, np.dot(xtrain, ridge.coef_) + ridge.intercept_)
    test_err2 = mean_squared_error(ytest, np.dot(xtest,ridge.coef_) + ridge.intercept_)

    #track min error for ridge
    if test_err2 < min_err[0]:
        min_err = (test_err2, split)

    errValuesRidge.append((train_err2, test_err2))

fig = plt.figure(figsize=(10,8))

xerrTrain = [e[0] for e in errValuesRegression]
xerrTest = [e[1] for e in errValuesRegression]

ax = fig.add_subplot(2,1,1)
ax.set_xlabel('split_size')
ax.set_title('Linear Regression')
ax.plot(splits, xerrTrain,label='Train' )
ax.plot(splits, xerrTest,label='Test' )
ax.legend()

xerrTrain2 = [e[0] for e in errValuesRidge]
xerrTest2 = [e[1] for e in errValuesRidge]

ax = fig.add_subplot(2,1,2)
ax.plot(splits, xerrTrain2,label='Train' )
ax.plot(splits, xerrTest2,label='Test' )
ax.set_xlabel('split_size')
ax.set_title('Ridge Regression')

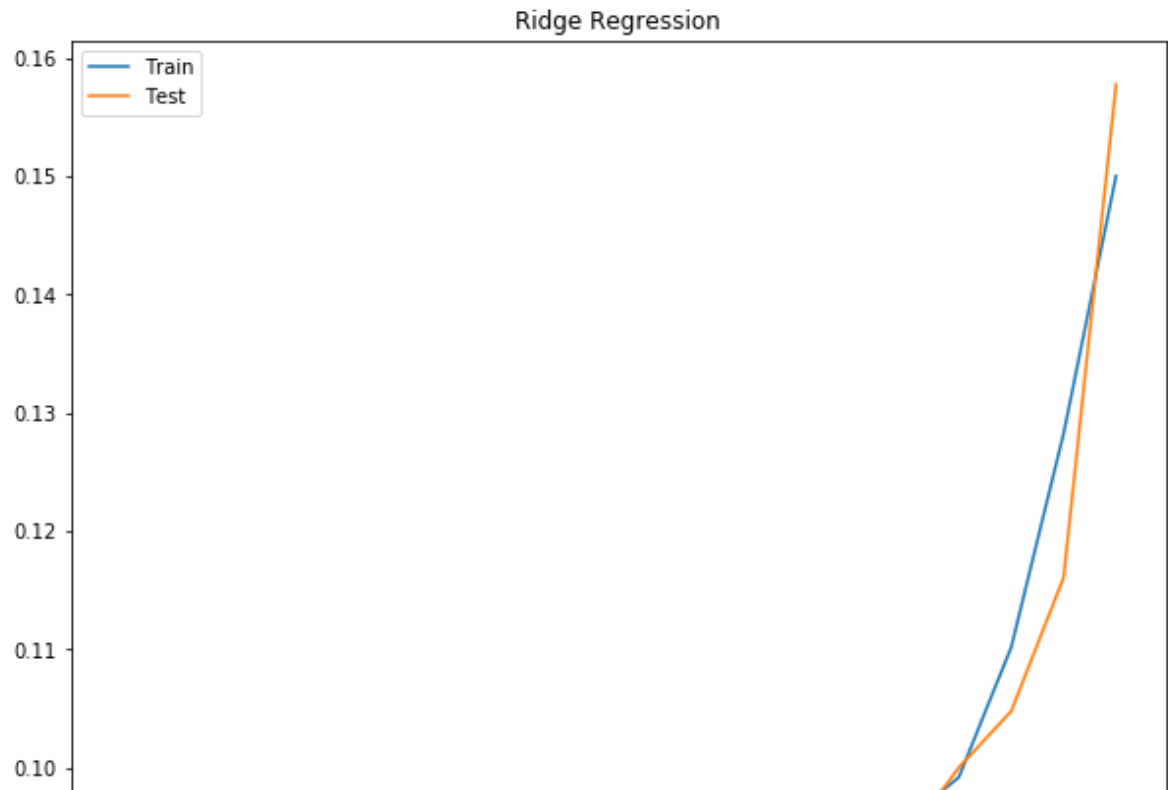
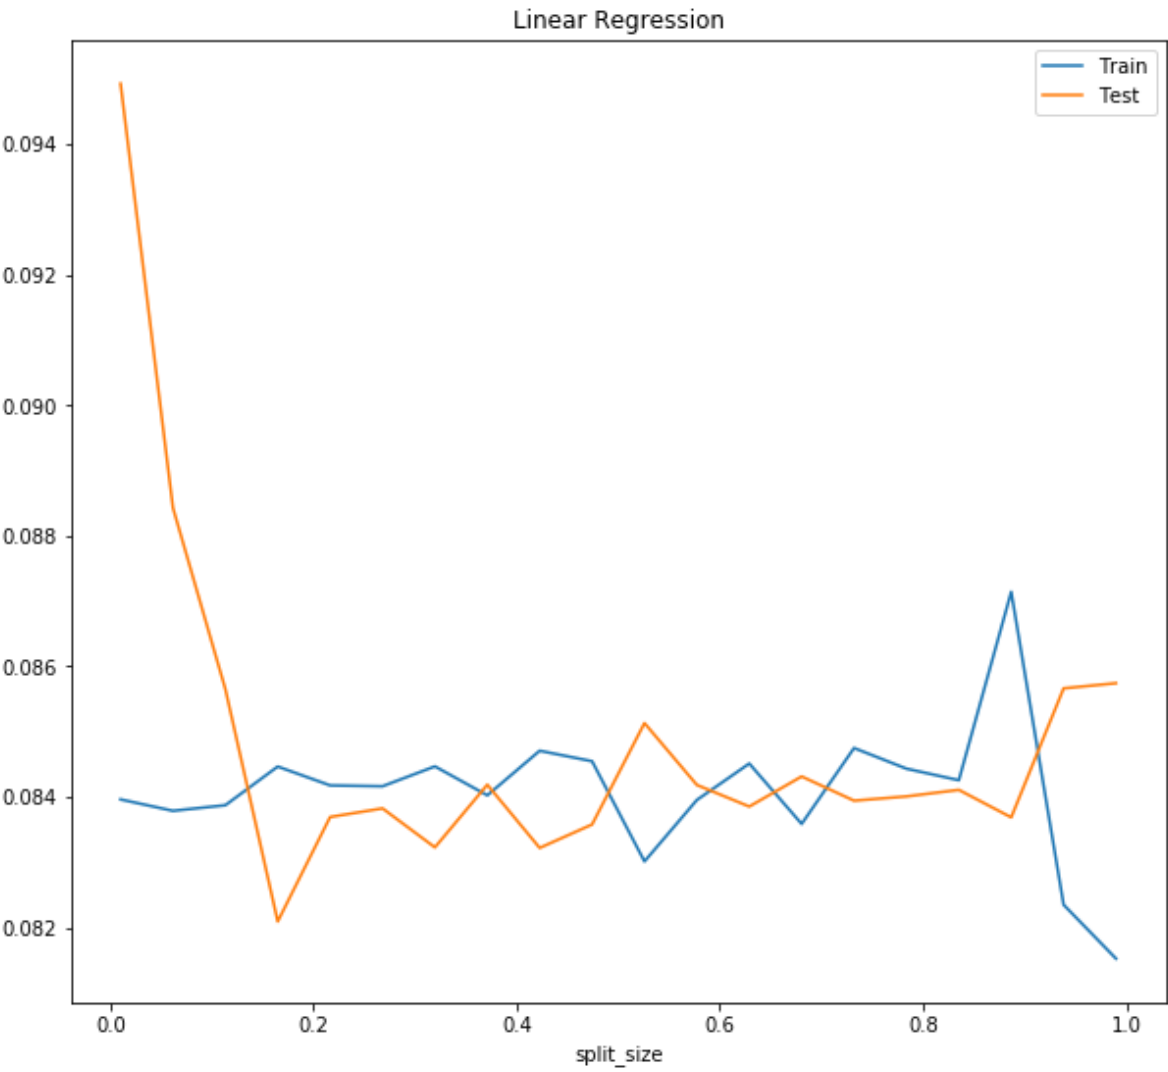
```


```
ax.legend()

fig.subplots_adjust(top=2.0)

print('min error {}'.format(min_err[0]))
print('optimal split at {}'.format(min_err[1]))
```

```
min error 0.08380034400734329  
optimal split at 0.16473684210526315
```





4. Chose an ideal split size based on the previous plot for Ridge.

Vary the Ridge parameter α from 0 to any value you'd like above 1. Plot the Train and Test error. Describe what you see based on the α parameter's stiffness.

```
In [35]: np.random.seed(1256)
alpha = np.linspace(0.0, 10000, 1000)
errValuesRidge = []

split = min_err[1]
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=split)


for a in alpha:
    ridge = Ridge(alpha=a,fit_intercept=True)
    ridge.fit(xtrain, ytrain)

    train_err = mean_squared_error(ytrain, np.dot(xtrain, ridge.coef_) + ridge
.intercept_)
    test_err = mean_squared_error(ytest, np.dot(xtest,ridge.coef_) + ridge.in
tercept_)

    errValuesRidge.append([train_err, test_err])

xerrTrain = [e[0] for e in errValuesRidge]
xerrTest = [e[1] for e in errValuesRidge]

plt.plot(alpha, xerrTrain,label='Train' )
plt.plot(alpha, xerrTest,label='Test' )
plt.xlabel('alpha')
plt.legend()
plt.show()
```



Bonus. Either: Generate data with a polynomial shape or use real data that you find on your own. Choose whatever regression model and process you'd like (Ridge, polynomial, etc.) and plot the Train-Test errors vs. any parameter your Model depends on (e.g. alpha, degree, etc.)