

## Import Appointment Data

```
In [1]: import pandas as pd
import numpy as np
#Load Data Section
data = pd.read_csv('./KaggleV2-May-2016.csv', parse_dates=[ 'ScheduledDay'
, 'AppointmentDay' ])
data.head()
```

Out[1]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbo
0	2.987250e+13	5642903	F	2016-04-29 18:38:08	2016-04-29	62	JARDIM PENHA
1	5.589978e+14	5642503	M	2016-04-29 16:08:27	2016-04-29	56	JARDIM PENHA
2	4.262962e+12	5642549	F	2016-04-29 16:19:04	2016-04-29	62	MATA DA
3	8.679512e+11	5642828	F	2016-04-29 17:29:31	2016-04-29	8	PONTAL CAMBUF
4	8.841186e+12	5642494	F	2016-04-29 16:07:23	2016-04-29	56	JARDIM PENHA

## MungeData

```

In [2]: #Munge Data
from sklearn.preprocessing import LabelEncoder
#Day of week, eg. Friday or Monday may be good predictor
data['ScheduleDayOfWeek'] = data.ScheduledDay.dt.dayofweek
data['AppointmentDayOfWeek'] = data.AppointmentDay.dt.dayofweek
data['SameDayAppt'] = np.where(data.ScheduledDay.dt.dayofweek == data.AppointmentDay.dt.dayofweek,1,0)
data['HoursUntilAppt'] = (data.AppointmentDay-data.ScheduledDay)
#convert yes, no to 0, 1
data['No-show-binary'] = np.where(data['No-show'] == 'Yes',1,0)
data['Gender-binary'] = np.where(data['Gender'] == 'Yes',1,0)

#convert neighbourhood into dummy columns
dummies = pd.get_dummies(data.Neighbourhood)
data[dummies.columns] = dummies
data.head()

```

Out[2]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbo
0	2.987250e+13	5642903	F	2016-04-29 18:38:08	2016-04-29	62	JARDIM PENHA
1	5.589978e+14	5642503	M	2016-04-29 16:08:27	2016-04-29	56	JARDIM PENHA
2	4.262962e+12	5642549	F	2016-04-29 16:19:04	2016-04-29	62	MATA DA
3	8.679512e+11	5642828	F	2016-04-29 17:29:31	2016-04-29	8	PONTAL CAMBU
4	8.841186e+12	5642494	F	2016-04-29 16:07:23	2016-04-29	56	JARDIM PENHA

5 rows × 101 columns

## Data Analysis

```

In [3]: num_columns = [col for col, dtype
    in zip(data.columns, data.dtypes) if dtype in ['float64', 'int64',
    'uint8'] and
    col not in ['No-show-binary', 'PatientId', 'AppointmentID'] ]

X = data[num_columns]
y = data['No-show-binary']

```

```
In [4]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split

logistic = LogisticRegression()
kn = KNeighborsClassifier()

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25
)

logistic.fit(x_train, y_train)
kn.fit(x_train, y_train)
```

```
Out[4]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowsk
i',
                             metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                             weights='uniform')
```

## Test Models

```
In [5]: from sklearn.metrics import accuracy_score, confusion_matrix, classifica
tion_report

#Logistic Regression
print('logistic accuracy score: ' + str(accuracy_score(y_test, logistic.
predict(x_test))))
print('logistic confusion matrix: ' + str(confusion_matrix(y_test, logis
tic.predict(x_test))))
print('logistic classification report')
print(classification_report(y_test, logistic.predict(x_test)))
```

```
logistic accuracy score: 0.796612623046
logistic confusion matrix: [[22006    10]
 [ 5610     6]]
logistic classification report
```

	precision	recall	f1-score	support
0	0.80	1.00	0.89	22016
1	0.38	0.00	0.00	5616
avg / total	0.71	0.80	0.71	27632

```
In [6]: #KNearest Neighbors
print('k-nearest neighbor accuracy score: ' + str(accuracy_score(y_test,
    kn.predict(x_test))))
print('k-nearest neighbor confusion matrix: ' + str(confusion_matrix(y_t
est, kn.predict(x_test))))
print('k-nearest classification report')
print(classification_report(y_test, kn.predict(x_test)))
```

```
k-nearest neighbor accuracy score: 0.772510133179
k-nearest neighbor confusion matrix: [[20397  1619]
 [ 4667   949]]
k-nearest classification report
```

	precision	recall	f1-score	support
0	0.81	0.93	0.87	22016
1	0.37	0.17	0.23	5616
avg / total	0.72	0.77	0.74	27632

## Try subset of data

### Remove neighbourhood dummy columns

```
In [7]: reduced_set = num_columns[0:11]
X = data[reduced_set]
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25
)

logistic.fit(x_train, y_train)
kn.fit(x_train, y_train)
```

```
Out[7]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowsk
i',
    metric_params=None, n_jobs=1, n_neighbors=5, p=2,
    weights='uniform')
```

```
In [8]: print('logistic accuracy score: ' + str(accuracy_score(y_test, logistic.
predict(x_test))))
print('logistic confusion matrix: ' + str(confusion_matrix(y_test, logistic.
predict(x_test))))
print('logistic classification report')
print(classification_report(y_test, logistic.predict(x_test)))
```

```
logistic accuracy score: 0.798204979734
logistic confusion matrix: [[22056      0]
 [ 5576      0]]
logistic classification report
```

	precision	recall	f1-score	support
0	0.80	1.00	0.89	22056
1	0.00	0.00	0.00	5576
avg / total	0.64	0.80	0.71	27632

```
/Users/jamescheever/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```

```
In [9]: print('k-nearest neighbor accuracy score: ' + str(accuracy_score(y_test,
kn.predict(x_test))))
print('k-nearest neighbor confusion matrix: ' + str(confusion_matrix(y_test,
kn.predict(x_test))))
print('k-nearest classification report')
print(classification_report(y_test, kn.predict(x_test)))
```

```
k-nearest neighbor accuracy score: 0.764837869137
k-nearest neighbor confusion matrix: [[20286 1770]
 [ 4728  848]]
k-nearest classification report
```

	precision	recall	f1-score	support
0	0.81	0.92	0.86	22056
1	0.32	0.15	0.21	5576
avg / total	0.71	0.76	0.73	27632

```
In [10]: X = data[['Age']]
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
logistic.fit(x_train, y_train)
print('logistic accuracy score: ' + str(accuracy_score(y_test, logistic.predict(x_test))))
print('logistic confusion matrix: ' + str(confusion_matrix(y_test, logistic.predict(x_test))))
print('logistic classification report')
print(classification_report(y_test, logistic.predict(x_test)))
```

```
logistic accuracy score: 0.797806890562
logistic confusion matrix: [[22045    0]
 [ 5587    0]]
logistic classification report
```

	precision	recall	f1-score	support
0	0.80	1.00	0.89	22045
1	0.00	0.00	0.00	5587
avg / total	0.64	0.80	0.71	27632

```
/Users/jamescheever/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:1135: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
```

```
'precision', 'predicted', average, warn_for)
```